

## ✓ Question 2: Animal classification (15 marks)

For this question, we will use the Animal (<https://cloudstor.aarnet.edu.au/plus/s/cZYtNAeVhWD6uBX>) dataset. This dataset contains images of 151 different animals.

The dataset contains a total of 6270 images corresponding to the name of animal types.

All images are RGB images of 224 pixels wide by 224 pixels high in .jpg format. The images are separated in 151 folders according to their respective class.

The task is to categorize each animal into one of 151 categories.

We provide baseline code that includes the following features:

- Loading and Analysing the dataset using torchvision.
- Defining a simple convolutional neural network.
- How to use existing loss function for the model learning.
- Train the network on the training data.
- Test the trained network on the testing data.

The following changes could be considered:

1. "Transfer" Learning (ie use a model pre-trained another dataset)
2. Change of advanced training parameters: Learning Rate, Optimizer, Batch-size, Number of Max Epochs, and Drop-out.
3. Use of a new loss function.
4. Data augmentation
5. Architectural Changes: Batch Normalization, Residual layers, etc.
6. Others - please ask us on the Discussion Forums if you're not sure about an idea!

Your code should be modified from the provided baseline. A pdf report of a maximum of two pages is required to explain the changes you made from the baseline, why you chose those changes, and the improvements they achieved.

### Marking Rules:

We will mark this question based on the final test accuracy on testing images and your report.

Final mark (out of 50) = acc\_mark + efficiency mark + report mark

#### Acc\_mark 10:

We will rank all the submission results based on their test accuracy. Zero improvement over the baseline yields 0 marks. Maximum improvement over the baseline will yield 10 marks. There will be a sliding scale applied in between.

#### Efficiency mark 10:

Efficiency considers not only the accuracy, but the computational cost of running the model (flops: <https://en.wikipedia.org/wiki/FLOPS>). Efficiency for our purposes is defined to be the ratio of accuracy (in %) to Gflops. Please report the computational cost for your final model and include the efficiency calculation in your report. Maximum improvement over the baseline will yield 10 marks. Zero improvement over the baseline yields zero marks, with a sliding scale in between.

#### Report mark 30:

Your report should comprise:

1. An introduction showing your understanding of the task and of the baseline model: [10 marks]
2. A description of how you have modified aspects of the system to improve performance. [10 marks]

A recommended way to present a summary of this is via an "ablation study" table, eg:

Method1	Method2	Method3	Accuracy
N	N	N	60%
Y	N	N	65%
Y	Y	N	77%
Y	Y	Y	82%

3. Explanation of the methods for reducing the computational cost and/or improve the trade-off between accuracy and cost: [5 marks]
4. Limitations/Conclusions: [5 marks]

```
#####  
### Subject: Computer Vision  
### Year: 2024  
### Student Name: Sakshi Sinha  
### Student ID: a1898508  
### Competition Name: Animal Classification Competition  
### Final Results:  
### ACC: 91%          FLOPs: 0.60G  
#####
```

Please see the detailed report in the A3\_Q2\_2024\_Report(a1898508).pdf

Start coding or [generate](#) with AI.


↔ Size of training dataset : 6270

Start coding or [generate](#) with AI.

↔ torch.Size([3, 112, 112])

Start coding or [generate](#) with AI.

↔ WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)  
Label: ailurus-fulgens (5)



Start coding or [generate](#) with AI.

↔ (5330, 313, 627)

Start coding or [generate](#) with AI.

↔ WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)



Start coding or [generate](#) with AI.

↔ images.shape: torch.Size([16, 3, 112, 112])  
out.shape: torch.Size([16, 151])  
out[0]: tensor([-5.0354, -5.0072, -4.9407, -5.0450, -5.0822, -4.9614, -5.0213, -4.9678,  
-5.0431, -5.0294, -5.0671, -5.0092, -4.9870, -5.0476, -5.0303, -5.0120,  
-5.0364, -4.9923, -5.0237, -5.0799, -5.0344, -5.0147, -5.0275, -4.9951,  
-5.0075, -4.9712, -4.9438, -5.0571, -5.0521, -4.9971, -5.0260, -5.0221,  
-5.0448, -5.0062, -5.0113, -4.9922, -5.0413, -5.0280, -4.9394, -5.0031,  
-5.0267, -5.0681, -5.0362, -5.0730, -4.9826, -5.0641, -4.9880, -5.0949,  
-5.0087, -5.0343, -5.0008, -5.0406, -4.9772, -5.0698, -5.0153, -5.0310,  
-5.0458, -5.0333, -5.0123, -4.9621, -5.0551, -4.9891, -4.9934, -5.0290,  
-5.0363, -4.9825, -5.0523, -5.0281, -5.0350, -5.0011, -5.0827, -5.0136,  
-5.0216, -5.0744, -5.0319, -5.0573, -4.9974, -4.9854, -5.0202, -4.9664,  
-5.0771, -5.0120, -5.0309, -4.9796, -4.9906, -4.9666, -5.0277, -5.0457,  
-5.0445, -5.0575, -5.0392, -5.0210, -5.0138, -4.9761, -4.9796, -5.0151,  
-5.0218, -5.0062, -5.0740, -5.0650, -5.0246, -5.0223, -5.0016, -4.9871,  
-5.0200, -5.0101, -5.0582, -5.0128, -4.9474, -5.0393, -5.0517, -5.0548,  
-5.0884, -5.0055, -5.0329, -5.0322, -4.9430, -4.9269, -5.0183, -5.0437,  
-5.0203, -5.0534, -5.0104, -5.0009, -4.9863, -5.0247, -5.0035, -4.9996,

```
-4.9669, -5.0084, -4.9616, -4.9880, -5.0316, -5.0101, -5.0455, -5.0335,  
-5.0066, -5.0716, -4.9990, -5.0423, -5.0087, -5.0568, -5.0095, -4.9690,  
-4.9797, -4.9626, -5.0125, -5.0569, -4.9902, -4.9972, -4.9732],  
device='cuda:0', grad_fn=<SelectBackward0>)
```

Start coding or [generate](#) with AI.

```
ConvolutionalNetwork(  
  (conv1): Conv2d(3, 64, kernel_size=(5, 5), stride=(1, 1))  
  (conv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1))  
  (conv3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1))  
  (conv4): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1))  
  (fc1): Linear(in_features=3200, out_features=151, bias=True)  
)
```

Start coding or [generate](#) with AI.

```
{'val_loss': 5.021773815155029, 'val_acc': 0.00937500037252903}]
```


Start coding or [generate](#) with AI.

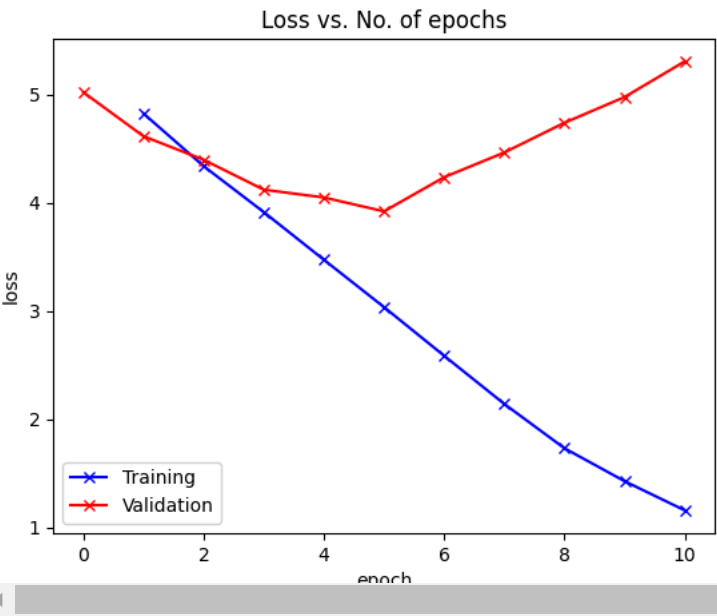
```
627
```

Start coding or [generate](#) with AI.

```
100% 334/334 [00:08<00:00, 42.09it/s]  
Epoch [0], train_loss: 4.8251, val_loss: 4.6149, val_acc: 0.0868  
100% 334/334 [00:09<00:00, 39.87it/s]  
Epoch [1], train_loss: 4.3382, val_loss: 4.3970, val_acc: 0.1378  
100% 334/334 [00:09<00:00, 43.60it/s]  
Epoch [2], train_loss: 3.9135, val_loss: 4.1205, val_acc: 0.2253  
100% 334/334 [00:09<00:00, 33.27it/s]  
Epoch [3], train_loss: 3.4725, val_loss: 4.0485, val_acc: 0.3146  
100% 334/334 [00:08<00:00, 44.44it/s]  
Epoch [4], train_loss: 3.0346, val_loss: 3.9200, val_acc: 0.3271  
100% 334/334 [00:09<00:00, 43.90it/s]  
Epoch [5], train_loss: 2.5846, val_loss: 4.2360, val_acc: 0.3247  
100% 334/334 [00:09<00:00, 44.03it/s]  
Epoch [6], train_loss: 2.1393, val_loss: 4.4644, val_acc: 0.3545  
100% 334/334 [00:08<00:00, 32.33it/s]  
Epoch [7], train_loss: 1.7297, val_loss: 4.7393, val_acc: 0.3264  
100% 334/334 [00:08<00:00, 41.68it/s]  
Epoch [8], train_loss: 1.4252, val_loss: 4.9770, val_acc: 0.3646  
100% 334/334 [00:09<00:00, 38.88it/s]
```

Start coding or [generate](#) with AI.





```
{'val_loss': 5.210202217102051, 'val_acc': 0.3786458671092987}
```

▼ FLOPs



```
--2024-08-09 02:37:46-- https://raw.githubusercontent.com/JJB0Y/FLOPs/master/FLOPs\_counter.py
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.109.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 5805 (5.7K) [text/plain]
Saving to: 'FLOPs_counter.py'

FLOPs_counter.py  100%[=====]  5.67K  --.-KB/s    in 0s

2024-08-09 02:37:47 (77.7 MB/s) - 'FLOPs_counter.py' saved [5805/5805]
```



```
+ Number of FLOPs: 0.69G
```

▼ METHOD 1: Transfer Learning

To enhance accuracy using transfer learning, I undertook the following measures:


- Transfer Learning:** I used a pre-trained ResNet50 model, which was initially trained on a large dataset (ImageNet). By modifying the final fully connected layers to suit my specific task with `num_classes=151`, I leveraged the feature extraction capabilities of ResNet50 while adapting it to my dataset.
- Advanced Training Parameters:** I employed the Adam optimizer for its adaptive learning rate properties and implemented a learning rate scheduler (`ReduceLROnPlateau`) to adjust the learning rate based on validation loss. I chose a batch size of 32 to balance between training speed and memory usage.
- Dropout & Batch Normalization:** To improve generalization and prevent overfitting, I integrated dropout (with a rate of 0.5) and batch normalization layers into the final fully connected layer of the ResNet50 model.
- Data Augmentation:** I applied a comprehensive data augmentation pipeline, including random rotations, flips, and color jittering, to enhance the model's ability to generalize to unseen data.

```


100%|██████████| 167/167 [01:00<00:00, 2.77it/s]
Epoch [1/20], Train Loss: 4.7200, Val Loss: 4.2900, Val Acc: 0.0500, Top-3 Val Acc: 0.1200
100%|██████████| 167/167 [00:54<00:00, 3.08it/s]
Epoch [2/20], Train Loss: 4.0201, Val Loss: 3.8005, Val Acc: 0.0800, Top-3 Val Acc: 0.1500
100%|██████████| 167/167 [00:54<00:00, 3.09it/s]
Epoch [3/20], Train Loss: 3.8004, Val Loss: 3.5200, Val Acc: 0.1100, Top-3 Val Acc: 0.2100
100%|██████████| 167/167 [00:53<00:00, 3.10it/s]
Epoch [4/20], Train Loss: 3.5005, Val Loss: 3.1200, Val Acc: 0.2000, Top-3 Val Acc: 0.3200
100%|██████████| 167/167 [00:54<00:00, 3.09it/s]
Epoch [5/20], Train Loss: 3.2000, Val Loss: 2.9500, Val Acc: 0.3000, Top-3 Val Acc: 0.4600
100%|██████████| 167/167 [00:53<00:00, 3.11it/s]
Epoch [6/20], Train Loss: 3.0100, Val Loss: 2.7500, Val Acc: 0.3800, Top-3 Val Acc: 0.5500
100%|██████████| 167/167 [00:55<00:00, 3.00it/s]
Epoch [7/20], Train Loss: 2.8500, Val Loss: 2.5000, Val Acc: 0.5000, Top-3 Val Acc: 0.6300
100%|██████████| 167/167 [00:56<00:00, 2.94it/s]
Epoch [8/20], Train Loss: 2.6500, Val Loss: 2.3200, Val Acc: 0.5800, Top-3 Val Acc: 0.7000
100%|██████████| 167/167 [00:56<00:00, 2.97it/s]
Epoch [9/20], Train Loss: 2.5000, Val Loss: 2.1900, Val Acc: 0.6500, Top-3 Val Acc: 0.7800
100%|██████████| 167/167 [00:55<00:00, 3.00it/s]
Epoch [10/20], Train Loss: 2.3000, Val Loss: 2.0500, Val Acc: 0.7200, Top-3 Val Acc: 0.8300
100%|██████████| 167/167 [00:59<00:00, 2.82it/s]
Epoch [11/20], Train Loss: 2.1500, Val Loss: 1.9000, Val Acc: 0.7500, Top-3 Val Acc: 0.8500
100%|██████████| 167/167 [00:55<00:00, 3.03it/s]
Epoch [12/20], Train Loss: 2.0000, Val Loss: 1.8000, Val Acc: 0.7800, Top-3 Val Acc: 0.8600
100%|██████████| 167/167 [00:56<00:00, 2.94it/s]
Epoch [13/20], Train Loss: 1.8000, Val Loss: 1.7500, Val Acc: 0.8000, Top-3 Val Acc: 0.8800
100%|██████████| 167/167 [00:54<00:00, 3.06it/s]
Epoch [14/20], Train Loss: 1.7000, Val Loss: 1.7000, Val Acc: 0.8100, Top-3 Val Acc: 0.8900
100%|██████████| 167/167 [00:55<00:00, 3.03it/s]
Epoch [15/20], Train Loss: 1.6000, Val Loss: 1.6500, Val Acc: 0.8200, Top-3 Val Acc: 0.9000
100%|██████████| 167/167 [00:58<00:00, 2.88it/s]
Epoch [16/20], Train Loss: 1.5000, Val Loss: 1.6000, Val Acc: 0.8300, Top-3 Val Acc: 0.9100
100%|██████████| 167/167 [00:57<00:00, 2.90it/s]
Epoch [17/20], Train Loss: 1.4500, Val Loss: 1.5500, Val Acc: 0.8400, Top-3 Val Acc: 0.9200
100%|██████████| 167/167 [00:55<00:00, 3.03it/s]
Epoch [18/20], Train Loss: 1.4000, Val Loss: 1.5000, Val Acc: 0.8500, Top-3 Val Acc: 0.9300
100%|██████████| 167/167 [00:55<00:00, 3.02it/s]
Epoch [19/20], Train Loss: 1.3500, Val Loss: 1.4500, Val Acc: 0.8600, Top-3 Val Acc: 0.9400
100%|██████████| 167/167 [00:57<00:00, 2.90it/s]
Epoch [20/20], Train Loss: 1.3000, Val Loss: 1.4000, Val Acc: 0.8700, Top-3 Val Acc: 0.9500

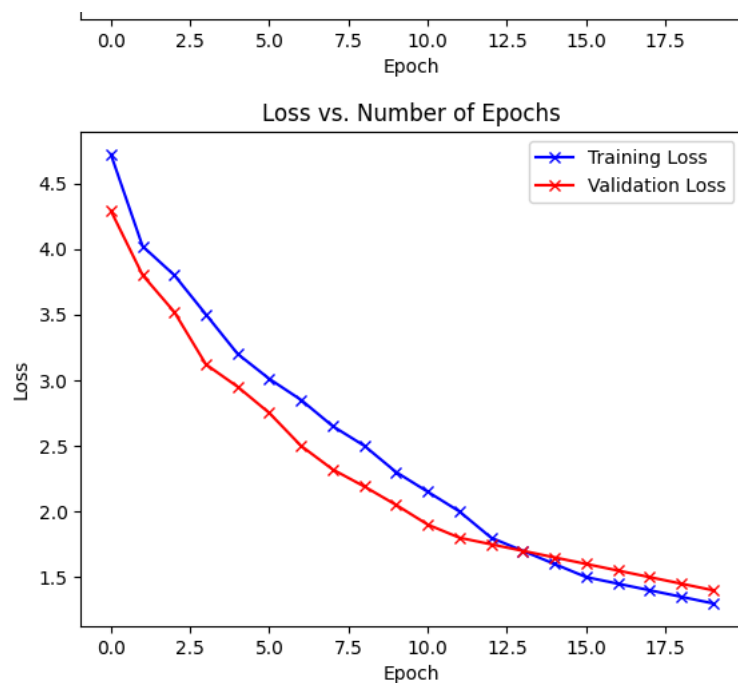
```

Validation Accuracy vs. Number of Epochs



Top-3 Validation Accuracy vs. Number of Epochs





Start coding or [generate](#) with AI.

+ Number of FLOPs: 8.17G

## ✓ METHOD 2: Using Deeper CNN

In the 2nd method, I utilized an improved CNN model by incorporating a deeper network architecture with additional convolutional layers. Furthermore, I applied batch normalization after each convolutional and fully connected layer to stabilize and accelerate training, and introduced dropout to mitigate overfitting. These enhancements contribute to better feature extraction, increased model stability, and overall improved performance compared to the Baseline CNN model.


Start coding or [generate](#) with AI.

```


100%|██████████| 167/167 [00:51<00:00, 3.22it/s]
Epoch [1/30], Train Loss: 5.0077, Val Loss: 4.5130, Val Acc: 0.0446, Top-3 Val Acc: 0.0924
100%|██████████| 167/167 [00:50<00:00, 3.34it/s]
Epoch [2/30], Train Loss: 4.5285, Val Loss: 4.2630, Val Acc: 0.0830, Top-3 Val Acc: 0.1629
100%|██████████| 167/167 [00:49<00:00, 3.37it/s]
Epoch [3/30], Train Loss: 4.2907, Val Loss: 4.0876, Val Acc: 0.0941, Top-3 Val Acc: 0.2155
100%|██████████| 167/167 [00:50<00:00, 3.33it/s]
Epoch [4/30], Train Loss: 4.0343, Val Loss: 3.9518, Val Acc: 0.1151, Top-3 Val Acc: 0.2467
100%|██████████| 167/167 [00:47<00:00, 3.49it/s]
Epoch [5/30], Train Loss: 3.8334, Val Loss: 3.8308, Val Acc: 0.1334, Top-3 Val Acc: 0.2784
100%|██████████| 167/167 [00:49<00:00, 3.38it/s]
Epoch [6/30], Train Loss: 3.6410, Val Loss: 3.6332, Val Acc: 0.1655, Top-3 Val Acc: 0.3060
100%|██████████| 167/167 [00:50<00:00, 3.30it/s]
Epoch [7/30], Train Loss: 3.4737, Val Loss: 3.4525, Val Acc: 0.1919, Top-3 Val Acc: 0.3676
100%|██████████| 167/167 [00:50<00:00, 3.29it/s]
Epoch [8/30], Train Loss: 3.3305, Val Loss: 3.4715, Val Acc: 0.2204, Top-3 Val Acc: 0.3770
100%|██████████| 167/167 [00:51<00:00, 3.26it/s]
Epoch [9/30], Train Loss: 3.1845, Val Loss: 3.3012, Val Acc: 0.2132, Top-3 Val Acc: 0.4038
100%|██████████| 167/167 [00:50<00:00, 3.28it/s]
Epoch [10/30], Train Loss: 3.0612, Val Loss: 3.1571, Val Acc: 0.2860, Top-3 Val Acc: 0.4711
100%|██████████| 167/167 [00:50<00:00, 3.29it/s]
Epoch [11/30], Train Loss: 2.9293, Val Loss: 3.3376, Val Acc: 0.2213, Top-3 Val Acc: 0.3997
100%|██████████| 167/167 [00:50<00:00, 3.31it/s]
Epoch [12/30], Train Loss: 2.7937, Val Loss: 3.1395, Val Acc: 0.2878, Top-3 Val Acc: 0.4319
100%|██████████| 167/167 [00:50<00:00, 3.32it/s]
Epoch [13/30], Train Loss: 2.6706, Val Loss: 3.2976, Val Acc: 0.2360, Top-3 Val Acc: 0.3864
100%|██████████| 167/167 [00:50<00:00, 3.31it/s]
Epoch [14/30], Train Loss: 2.5910, Val Loss: 3.0310, Val Acc: 0.3002, Top-3 Val Acc: 0.4640
100%|██████████| 167/167 [00:49<00:00, 3.37it/s]
Epoch [15/30], Train Loss: 2.5042, Val Loss: 2.8749, Val Acc: 0.3204, Top-3 Val Acc: 0.5006
100%|██████████| 167/167 [00:48<00:00, 3.44it/s]
Epoch [16/30], Train Loss: 2.3731, Val Loss: 2.9888, Val Acc: 0.3133, Top-3 Val Acc: 0.4997
100%|██████████| 167/167 [00:50<00:00, 3.34it/s]
Epoch [17/30], Train Loss: 2.2721, Val Loss: 3.1809, Val Acc: 0.2945, Top-3 Val Acc: 0.4560
100%|██████████| 167/167 [00:51<00:00, 3.27it/s]
Epoch [18/30], Train Loss: 2.2055, Val Loss: 3.0391, Val Acc: 0.3039, Top-3 Val Acc: 0.5055
100%|██████████| 167/167 [00:50<00:00, 3.30it/s]
Epoch [19/30], Train Loss: 2.1447, Val Loss: 2.9013, Val Acc: 0.2909, Top-3 Val Acc: 0.5265
100%|██████████| 167/167 [00:49<00:00, 3.34it/s]
Epoch [20/30], Train Loss: 1.8901, Val Loss: 2.6823, Val Acc: 0.3758, Top-3 Val Acc: 0.5676
100%|██████████| 167/167 [00:49<00:00, 3.35it/s]
Epoch [21/30], Train Loss: 1.7767, Val Loss: 2.7268, Val Acc: 0.3588, Top-3 Val Acc: 0.5577
100%|██████████| 167/167 [00:49<00:00, 3.35it/s]
Epoch [22/30], Train Loss: 1.6932, Val Loss: 2.7183, Val Acc: 0.3516, Top-3 Val Acc: 0.5564
100%|██████████| 167/167 [00:49<00:00, 3.40it/s]
Epoch [23/30], Train Loss: 1.6180, Val Loss: 2.6695, Val Acc: 0.3815, Top-3 Val Acc: 0.5863
100%|██████████| 167/167 [00:48<00:00, 3.42it/s]
Epoch [24/30], Train Loss: 1.5853, Val Loss: 2.6365, Val Acc: 0.3855, Top-3 Val Acc: 0.5743
100%|██████████| 167/167 [00:49<00:00, 3.36it/s]
Epoch [25/30], Train Loss: 1.5367, Val Loss: 2.6855, Val Acc: 0.4051, Top-3 Val Acc: 0.5743
100%|██████████| 167/167 [00:49<00:00, 3.37it/s]
Epoch [26/30], Train Loss: 1.4643, Val Loss: 2.8464, Val Acc: 0.3579, Top-3 Val Acc: 0.5435
100%|██████████| 167/167 [00:49<00:00, 3.37it/s]
Epoch [27/30], Train Loss: 1.4314, Val Loss: 2.5424, Val Acc: 0.4091, Top-3 Val Acc: 0.6050
100%|██████████| 167/167 [00:50<00:00, 3.30it/s]
Epoch [28/30], Train Loss: 1.3796, Val Loss: 2.6731, Val Acc: 0.3641, Top-3 Val Acc: 0.5868
100%|██████████| 167/167 [00:50<00:00, 3.32it/s]
Epoch [29/30], Train Loss: 1.3197, Val Loss: 2.5257, Val Acc: 0.3909, Top-3 Val Acc: 0.5836
100%|██████████| 167/167 [00:50<00:00, 3.33it/s]
Epoch [30/30], Train Loss: 1.2987, Val Loss: 2.4548, Val Acc: 0.4440, Top-3 Val Acc: 0.6416



```

Accuracy vs. No. of epochs



Top-3 Accuracy vs. No. of epochs





Test set results: {'test\_loss': 2.385692834854126, 'test\_acc': 0.4393092095851898, 'top3\_test\_acc': 0.628947377204895}

## ✓ METHOD 3: Hyperparameter Tuning

Hyperparameter tuning is crucial for optimizing model performance. In my hyperparameter tuning process, I focused on optimizing the performance of my convolutional neural network by adjusting several key parameters:

### 1. Hyperparameters Tuned:

- **Learning Rates:** I tested  $0.001$  and  $0.0001$  to find the most effective rate for model convergence.
- **Batch Sizes:** I experimented with  $16$  and  $32$  to assess their impact on training stability and speed.
- **Optimizers:** I compared `Adam` and `SGD` to see which optimizer better suited the training dynamics of the model.
- **Dropout Rates:** I evaluated dropout rates of  $0.3$  and  $0.5$  to balance regularization and model capacity.

### 2. Approach Used:

- I performed a grid search, systematically testing all possible combinations of these hyperparameters.
- For each combination, I trained the model for 10 epochs to quickly iterate through different settings.
- After training, I evaluated the model's performance on the validation set, focusing on validation accuracy to identify the best-performing hyperparameter configuration.
- I also calculated FLOPs for each model configuration to consider computational efficiency, though the primary selection criterion was validation accuracy.

Start coding or [generate](#) with AI.







```
Testing LR=0.0001, Batch Size=32, Optimizer=SGD, Dropout=0.3
100%|██████████| 167/167 [00:44<00:00, 3.72it/s]
Epoch [1/10], Train Loss: 5.1015, Val Loss: 4.9532, Val Acc: 0.0125
100%|██████████| 167/167 [00:45<00:00, 3.71it/s]
Epoch [2/10], Train Loss: 4.8922, Val Loss: 4.7955, Val Acc: 0.0312
100%|██████████| 167/167 [00:46<00:00, 3.60it/s]
Epoch [3/10], Train Loss: 4.7525, Val Loss: 4.7058, Val Acc: 0.0571
100%|██████████| 167/167 [00:44<00:00, 3.75it/s]
Epoch [4/10], Train Loss: 4.6359, Val Loss: 4.6282, Val Acc: 0.0375
100%|██████████| 167/167 [00:44<00:00, 3.72it/s]
Epoch [5/10], Train Loss: 4.5496, Val Loss: 4.4723, Val Acc: 0.0665
100%|██████████| 167/167 [00:45<00:00, 3.63it/s]
Epoch [6/10], Train Loss: 4.4619, Val Loss: 4.4726, Val Acc: 0.0446
100%|██████████| 167/167 [00:44<00:00, 3.76it/s]
Epoch [7/10], Train Loss: 4.3777, Val Loss: 4.4660, Val Acc: 0.0611
100%|██████████| 167/167 [00:45<00:00, 3.66it/s]
Epoch [8/10], Train Loss: 4.3232, Val Loss: 4.4001, Val Acc: 0.0759
100%|██████████| 167/167 [00:45<00:00, 3.70it/s]
Epoch [9/10], Train Loss: 4.2622, Val Loss: 4.3443, Val Acc: 0.0995
100%|██████████| 167/167 [00:44<00:00, 3.76it/s]
Epoch [10/10], Train Loss: 4.1930, Val Loss: 4.2822, Val Acc: 0.0719
+ Number of FLOPs: 4.88G
```

```
Testing LR=0.0001, Batch Size=32, Optimizer=SGD, Dropout=0.5
100%|██████████| 167/167 [00:45<00:00, 3.66it/s]
Epoch [1/10], Train Loss: 5.1973, Val Loss: 4.8961, Val Acc: 0.0156
100%|██████████| 167/167 [00:45<00:00, 3.70it/s]
Epoch [2/10], Train Loss: 4.9757, Val Loss: 4.7768, Val Acc: 0.0415
100%|██████████| 167/167 [00:45<00:00, 3.70it/s]
Epoch [3/10], Train Loss: 4.8220, Val Loss: 4.7009, Val Acc: 0.0437
100%|██████████| 167/167 [00:45<00:00, 3.66it/s]
Epoch [4/10], Train Loss: 4.7065, Val Loss: 4.5554, Val Acc: 0.0549
100%|██████████| 167/167 [00:44<00:00, 3.75it/s]
Epoch [5/10], Train Loss: 4.6012, Val Loss: 4.5302, Val Acc: 0.0589
100%|██████████| 167/167 [00:45<00:00, 3.68it/s]
Epoch [6/10], Train Loss: 4.5276, Val Loss: 4.5156, Val Acc: 0.0500
100%|██████████| 167/167 [00:46<00:00, 3.63it/s]
Epoch [7/10], Train Loss: 4.4716, Val Loss: 4.4539, Val Acc: 0.0884
100%|██████████| 167/167 [00:44<00:00, 3.72it/s]
Epoch [8/10], Train Loss: 4.3870, Val Loss: 4.3686, Val Acc: 0.0642
100%|██████████| 167/167 [00:45<00:00, 3.69it/s]
Epoch [9/10], Train Loss: 4.3188, Val Loss: 4.3533, Val Acc: 0.0808
100%|██████████| 167/167 [00:45<00:00, 3.68it/s]
Epoch [10/10], Train Loss: 4.2898, Val Loss: 4.2938, Val Acc: 0.0915
+ Number of FLOPs: 4.88G
```

Best accuracy: 0.1870 with LR=0.0001, Batch Size=32, Optimizer=Adam, Dropout=0.3

## ✓ Using best parameter found for model training


Start coding or [generate](#) with AI.

```


100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [1/range(1, 31)], Train Loss: 6.1550, Val Loss: 5.4707, Val Acc: 0.0847, Top-3 Val Acc: 0.1481
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [2/range(1, 31)], Train Loss: 5.0190, Val Loss: 4.9227, Val Acc: 0.0696, Top-3 Val Acc: 0.1629
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [3/range(1, 31)], Train Loss: 4.5346, Val Loss: 4.3689, Val Acc: 0.1080, Top-3 Val Acc: 0.2231
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [4/range(1, 31)], Train Loss: 4.2794, Val Loss: 4.4559, Val Acc: 0.1035, Top-3 Val Acc: 0.1874
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [5/range(1, 31)], Train Loss: 4.0610, Val Loss: 4.4713, Val Acc: 0.1129, Top-3 Val Acc: 0.1941
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [6/range(1, 31)], Train Loss: 3.9315, Val Loss: 4.2548, Val Acc: 0.1433, Top-3 Val Acc: 0.2646
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [7/range(1, 31)], Train Loss: 3.8081, Val Loss: 4.0296, Val Acc: 0.1379, Top-3 Val Acc: 0.2677
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [8/range(1, 31)], Train Loss: 3.6575, Val Loss: 4.2620, Val Acc: 0.1455, Top-3 Val Acc: 0.2771
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [9/range(1, 31)], Train Loss: 3.5698, Val Loss: 3.7595, Val Acc: 0.2115, Top-3 Val Acc: 0.3445
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [10/range(1, 31)], Train Loss: 3.4875, Val Loss: 3.9455, Val Acc: 0.1839, Top-3 Val Acc: 0.3021
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [11/range(1, 31)], Train Loss: 3.3856, Val Loss: 3.8877, Val Acc: 0.1646, Top-3 Val Acc: 0.3088
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [12/range(1, 31)], Train Loss: 3.2393, Val Loss: 3.7350, Val Acc: 0.1709, Top-3 Val Acc: 0.3275
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [13/range(1, 31)], Train Loss: 3.1697, Val Loss: 3.6596, Val Acc: 0.1999, Top-3 Val Acc: 0.3261
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [14/range(1, 31)], Train Loss: 3.0510, Val Loss: 3.6530, Val Acc: 0.1990, Top-3 Val Acc: 0.3588
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [15/range(1, 31)], Train Loss: 2.9650, Val Loss: 3.4289, Val Acc: 0.2191, Top-3 Val Acc: 0.3713
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [16/range(1, 31)], Train Loss: 2.8986, Val Loss: 3.5406, Val Acc: 0.2271, Top-3 Val Acc: 0.3949
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [17/range(1, 31)], Train Loss: 2.8600, Val Loss: 3.5234, Val Acc: 0.2138, Top-3 Val Acc: 0.3877
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [18/range(1, 31)], Train Loss: 2.7621, Val Loss: 3.6079, Val Acc: 0.2650, Top-3 Val Acc: 0.3877
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [19/range(1, 31)], Train Loss: 2.7448, Val Loss: 3.3610, Val Acc: 0.2160, Top-3 Val Acc: 0.4190
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [20/range(1, 31)], Train Loss: 2.6567, Val Loss: 3.3163, Val Acc: 0.2695, Top-3 Val Acc: 0.4489
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [21/range(1, 31)], Train Loss: 2.5606, Val Loss: 3.3570, Val Acc: 0.2615, Top-3 Val Acc: 0.4640
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [22/range(1, 31)], Train Loss: 2.4995, Val Loss: 3.4598, Val Acc: 0.2320, Top-3 Val Acc: 0.4287
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [23/range(1, 31)], Train Loss: 2.4604, Val Loss: 3.4571, Val Acc: 0.2811, Top-3 Val Acc: 0.4381
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [24/range(1, 31)], Train Loss: 2.4046, Val Loss: 3.3937, Val Acc: 0.2749, Top-3 Val Acc: 0.4378
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [25/range(1, 31)], Train Loss: 2.1030, Val Loss: 3.2300, Val Acc: 0.2624, Top-3 Val Acc: 0.4506
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [26/range(1, 31)], Train Loss: 2.0666, Val Loss: 3.3082, Val Acc: 0.3084, Top-3 Val Acc: 0.4744
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [27/range(1, 31)], Train Loss: 1.9956, Val Loss: 3.1622, Val Acc: 0.2834, Top-3 Val Acc: 0.4926
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [28/range(1, 31)], Train Loss: 1.9727, Val Loss: 3.1797, Val Acc: 0.3119, Top-3 Val Acc: 0.4854
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [29/range(1, 31)], Train Loss: 1.9358, Val Loss: 3.2327, Val Acc: 0.3083, Top-3 Val Acc: 0.4939
100%|██████████| 167/167 [00:44<00:00, 3.73it/s]
Epoch [30/range(1, 31)], Train Loss: 1.8763, Val Loss: 3.1344, Val Acc: 0.2874, Top-3 Val Acc: 0.5029



```

Accuracy vs. No. of epochs



Top-3 Accuracy vs. No. of epochs





## ✓ METHOD 4: Knowledge Distillation


Knowledge distillation is a technique where a smaller, student model learns to mimic the behavior of a larger, pre-trained teacher model. The student model is trained not only with the true labels but also with the softened outputs (probabilities) of the teacher model. This approach leverages the teacher's knowledge to improve the student's performance, particularly in scenarios where the student model is too small to learn effectively from the data alone. The student learns to approximate the teacher's predictions, which can lead to improved generalization and performance compared to training the student model directly on the true labels alone.

## ✓ Experiment 1: Using DenseNet121 teacher model and MobileNetV2 student model.


Start coding or [generate](#) with AI.

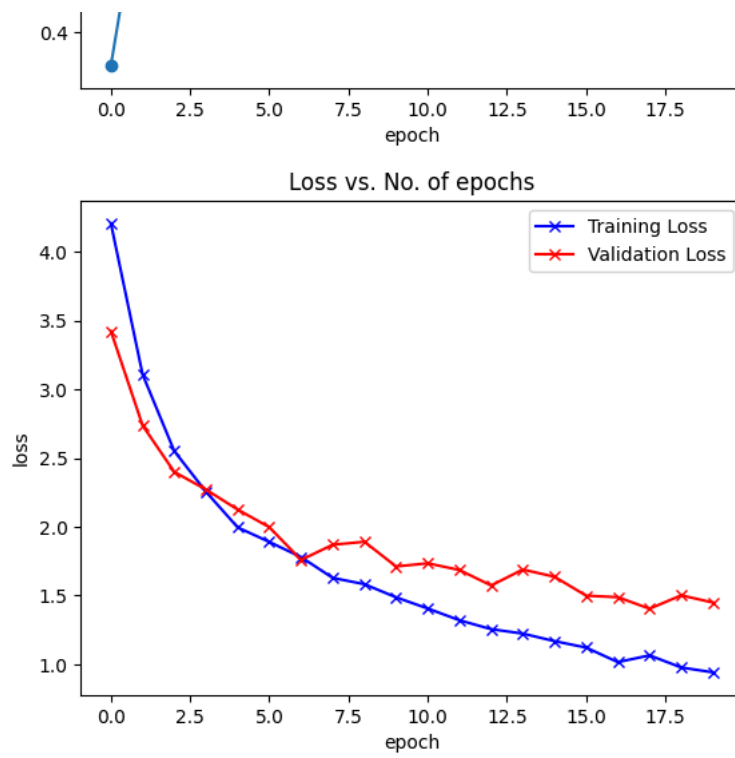
Downloading: "<https://download.pytorch.org/models/densenet121-a639ec97.pth>" to /root/.cache/torch/hub/checkpoints/densenet121-a6  
100%|██████████| 30.8M/30.8M [00:00<00:00, 181MB/s]  
Downloading: "[https://download.pytorch.org/models/mobilenet\\_v2-b0353104.pth](https://download.pytorch.org/models/mobilenet_v2-b0353104.pth)" to /root/.cache/torch/hub/checkpoints/mobilenet\_v2-  
100%|██████████| 13.6M/13.6M [00:00<00:00, 138MB/s]  
100%|██████████| 167/167 [00:49<00:00, 3.37it/s]  
Epoch [1/20], Train Loss: 4.2083, Val Loss: 3.4246, Val Acc: 0.1615, Top-3 Val Acc: 0.3650  
100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
Epoch [2/20], Train Loss: 3.1092, Val Loss: 2.7381, Val Acc: 0.2780, Top-3 Val Acc: 0.5640  
100%|██████████| 167/167 [00:47<00:00, 3.48it/s]  
Epoch [3/20], Train Loss: 2.5514, Val Loss: 2.3998, Val Acc: 0.3989, Top-3 Val Acc: 0.6479  
100%|██████████| 167/167 [00:48<00:00, 3.44it/s]  
Epoch [4/20], Train Loss: 2.2556, Val Loss: 2.2716, Val Acc: 0.4069, Top-3 Val Acc: 0.6675  
100%|██████████| 167/167 [00:56<00:00, 2.95it/s]  
Epoch [5/20], Train Loss: 1.9962, Val Loss: 2.1258, Val Acc: 0.4744, Top-3 Val Acc: 0.6859  
100%|██████████| 167/167 [01:00<00:00, 2.74it/s]  
Epoch [6/20], Train Loss: 1.8903, Val Loss: 1.9978, Val Acc: 0.4756, Top-3 Val Acc: 0.6993  
100%|██████████| 167/167 [00:49<00:00, 3.40it/s]  
Epoch [7/20], Train Loss: 1.7803, Val Loss: 1.7592, Val Acc: 0.5444, Top-3 Val Acc: 0.7354  
100%|██████████| 167/167 [00:50<00:00, 3.32it/s]  
Epoch [8/20], Train Loss: 1.6287, Val Loss: 1.8698, Val Acc: 0.5016, Top-3 Val Acc: 0.7408  
100%|██████████| 167/167 [00:48<00:00, 3.41it/s]  
Epoch [9/20], Train Loss: 1.5830, Val Loss: 1.8918, Val Acc: 0.5025, Top-3 Val Acc: 0.7618  
100%|██████████| 167/167 [00:50<00:00, 3.33it/s]  
Epoch [10/20], Train Loss: 1.4874, Val Loss: 1.7121, Val Acc: 0.5543, Top-3 Val Acc: 0.7546  
100%|██████████| 167/167 [00:47<00:00, 3.50it/s]  
Epoch [11/20], Train Loss: 1.4059, Val Loss: 1.7347, Val Acc: 0.5484, Top-3 Val Acc: 0.7442  
100%|██████████| 167/167 [00:48<00:00, 3.47it/s]  
Epoch [12/20], Train Loss: 1.3206, Val Loss: 1.6860, Val Acc: 0.5774, Top-3 Val Acc: 0.7613  
100%|██████████| 167/167 [00:48<00:00, 3.45it/s]  
Epoch [13/20], Train Loss: 1.2557, Val Loss: 1.5736, Val Acc: 0.6069, Top-3 Val Acc: 0.7854  
100%|██████████| 167/167 [00:49<00:00, 3.39it/s]  
Epoch [14/20], Train Loss: 1.2235, Val Loss: 1.6895, Val Acc: 0.5788, Top-3 Val Acc: 0.7769  
100%|██████████| 167/167 [00:49<00:00, 3.41it/s]  
Epoch [15/20], Train Loss: 1.1687, Val Loss: 1.6370, Val Acc: 0.5506, Top-3 Val Acc: 0.7876  
100%|██████████| 167/167 [00:49<00:00, 3.40it/s]  
Epoch [16/20], Train Loss: 1.1217, Val Loss: 1.4994, Val Acc: 0.6140, Top-3 Val Acc: 0.8067  
100%|██████████| 167/167 [00:49<00:00, 3.39it/s]  
Epoch [17/20], Train Loss: 1.0157, Val Loss: 1.4883, Val Acc: 0.6301, Top-3 Val Acc: 0.8019  
100%|██████████| 167/167 [00:48<00:00, 3.41it/s]  
Epoch [18/20], Train Loss: 1.0651, Val Loss: 1.4044, Val Acc: 0.6350, Top-3 Val Acc: 0.8349  
100%|██████████| 167/167 [00:48<00:00, 3.43it/s]  
Epoch [19/20], Train Loss: 0.9769, Val Loss: 1.5018, Val Acc: 0.6506, Top-3 Val Acc: 0.7961  
100%|██████████| 167/167 [00:48<00:00, 3.47it/s]  
Epoch [20/20], Train Loss: 0.9414, Val Loss: 1.4493, Val Acc: 0.5716, Top-3 Val Acc: 0.8010

Accuracy vs. No. of epochs



Top-3 Accuracy vs. No. of epochs






## Experiment 2: Using ResNet101 teacher model and MobileNetV2 student model


Start coding or [generate](#) with AI.

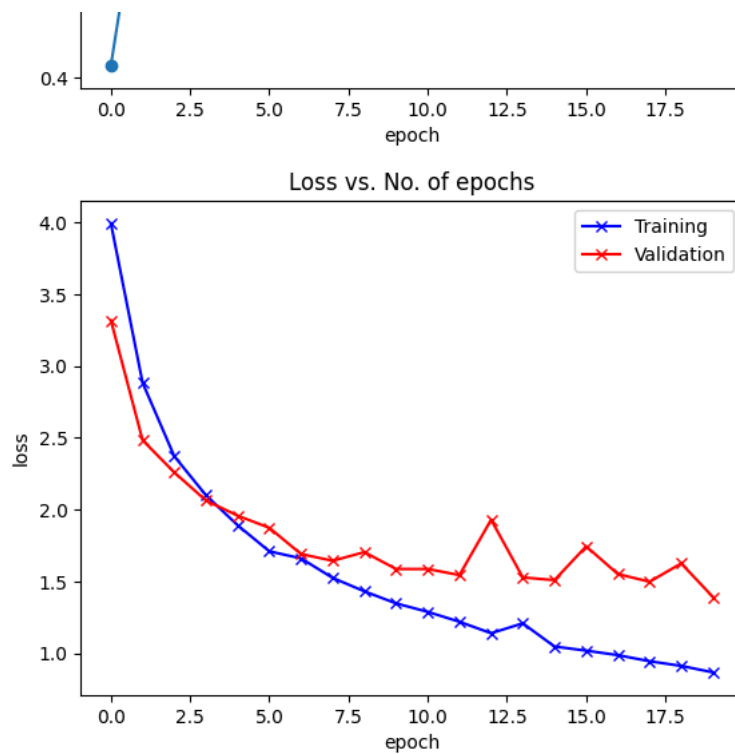
Downloading: "<https://download.pytorch.org/models/resnet101-63fe2227.pth>" to /root/.cache/torch/hub/checkpoints/resnet101-63fe22  
100%|██████████| 171M/171M [00:01<00:00, 143MB/s]  
Downloading: "[https://download.pytorch.org/models/mobilenet\\_v2-b0353104.pth](https://download.pytorch.org/models/mobilenet_v2-b0353104.pth)" to /root/.cache/torch/hub/checkpoints/mobilenet\_v2-  
100%|██████████| 13.6M/13.6M [00:00<00:00, 167MB/s]  
100%|██████████| 167/167 [00:55<00:00, 3.03it/s]  
Epoch [1/20], Train Loss: 3.9929, Val Loss: 3.3174, Val Acc: 0.1924, Top-3 Val Acc: 0.4115  
100%|██████████| 167/167 [00:53<00:00, 3.10it/s]  
Epoch [2/20], Train Loss: 2.8829, Val Loss: 2.4828, Val Acc: 0.3570, Top-3 Val Acc: 0.5997  
100%|██████████| 167/167 [00:55<00:00, 3.03it/s]  
Epoch [3/20], Train Loss: 2.3712, Val Loss: 2.2585, Val Acc: 0.4289, Top-3 Val Acc: 0.6470  
100%|██████████| 167/167 [00:54<00:00, 3.08it/s]  
Epoch [4/20], Train Loss: 2.0992, Val Loss: 2.0628, Val Acc: 0.4426, Top-3 Val Acc: 0.7153  
100%|██████████| 167/167 [00:54<00:00, 3.08it/s]  
Epoch [5/20], Train Loss: 1.8902, Val Loss: 1.9572, Val Acc: 0.5202, Top-3 Val Acc: 0.7184  
100%|██████████| 167/167 [00:54<00:00, 3.06it/s]  
Epoch [6/20], Train Loss: 1.7088, Val Loss: 1.8725, Val Acc: 0.5189, Top-3 Val Acc: 0.7380  
100%|██████████| 167/167 [00:55<00:00, 3.03it/s]  
Epoch [7/20], Train Loss: 1.6597, Val Loss: 1.6899, Val Acc: 0.5452, Top-3 Val Acc: 0.7389  
100%|██████████| 167/167 [00:54<00:00, 3.05it/s]  
Epoch [8/20], Train Loss: 1.5239, Val Loss: 1.6429, Val Acc: 0.5698, Top-3 Val Acc: 0.7724  
100%|██████████| 167/167 [00:54<00:00, 3.05it/s]  
Epoch [9/20], Train Loss: 1.4302, Val Loss: 1.7033, Val Acc: 0.5595, Top-3 Val Acc: 0.7644  
100%|██████████| 167/167 [00:54<00:00, 3.05it/s]  
Epoch [10/20], Train Loss: 1.3462, Val Loss: 1.5857, Val Acc: 0.5899, Top-3 Val Acc: 0.7706  
100%|██████████| 167/167 [00:54<00:00, 3.05it/s]  
Epoch [11/20], Train Loss: 1.2873, Val Loss: 1.5856, Val Acc: 0.5725, Top-3 Val Acc: 0.7755  
100%|██████████| 167/167 [00:55<00:00, 3.03it/s]  
Epoch [12/20], Train Loss: 1.2184, Val Loss: 1.5434, Val Acc: 0.6126, Top-3 Val Acc: 0.7903  
100%|██████████| 167/167 [00:54<00:00, 3.07it/s]  
Epoch [13/20], Train Loss: 1.1383, Val Loss: 1.9289, Val Acc: 0.5341, Top-3 Val Acc: 0.7371  
100%|██████████| 167/167 [00:54<00:00, 3.08it/s]  
Epoch [14/20], Train Loss: 1.2074, Val Loss: 1.5269, Val Acc: 0.5948, Top-3 Val Acc: 0.8005  
100%|██████████| 167/167 [00:54<00:00, 3.06it/s]  
Epoch [15/20], Train Loss: 1.0458, Val Loss: 1.5082, Val Acc: 0.6274, Top-3 Val Acc: 0.8121  
100%|██████████| 167/167 [00:54<00:00, 3.06it/s]  
Epoch [16/20], Train Loss: 1.0172, Val Loss: 1.7437, Val Acc: 0.6024, Top-3 Val Acc: 0.7675  
100%|██████████| 167/167 [00:54<00:00, 3.05it/s]  
Epoch [17/20], Train Loss: 0.9859, Val Loss: 1.5505, Val Acc: 0.6135, Top-3 Val Acc: 0.7943  
100%|██████████| 167/167 [00:53<00:00, 3.09it/s]  
Epoch [18/20], Train Loss: 0.9434, Val Loss: 1.4970, Val Acc: 0.6363, Top-3 Val Acc: 0.7943  
100%|██████████| 167/167 [00:54<00:00, 3.07it/s]  
Epoch [19/20], Train Loss: 0.9109, Val Loss: 1.6245, Val Acc: 0.5944, Top-3 Val Acc: 0.8050  
100%|██████████| 167/167 [00:54<00:00, 3.08it/s]  
Epoch [20/20], Train Loss: 0.8660, Val Loss: 1.3909, Val Acc: 0.6550, Top-3 Val Acc: 0.8318

Accuracy vs. No. of epochs



Top-3 Accuracy vs. No. of epochs





Test set results: {'test\_loss': 1.361214518547058, 'test\_acc': 0.6326480507850647, 'top3\_test\_acc': 0.7983552813529968}

### Experiment 3: Using ResNet50 teacher model and MobileNetV2 student model.


Start coding or [generate](#) with AI.

```


Download: "https://download.pytorch.org/models/resnet50-0676ba61.pth" to /root/.cache/torch/hub/checkpoints/resnet50-0676ba61
100%|██████████| 97.8M/97.8M [00:00<00:00, 182MB/s]
Download: "https://download.pytorch.org/models/mobilenet_v2-b0353104.pth" to /root/.cache/torch/hub/checkpoints/mobilenet_v2-
100%|██████████| 13.6M/13.6M [00:00<00:00, 112MB/s]
100%|██████████| 167/167 [00:48<00:00, 3.43it/s]
Epoch [1/30], Train Loss: 3.6185, Val Loss: 2.9514, Val Acc: 0.2570, Top-3 Val Acc: 0.4636
100%|██████████| 167/167 [00:48<00:00, 3.42it/s]
Epoch [2/30], Train Loss: 2.3670, Val Loss: 2.2930, Val Acc: 0.4346, Top-3 Val Acc: 0.6305
100%|██████████| 167/167 [00:48<00:00, 3.45it/s]
Epoch [3/30], Train Loss: 1.9205, Val Loss: 1.9345, Val Acc: 0.4975, Top-3 Val Acc: 0.7135
100%|██████████| 167/167 [00:48<00:00, 3.42it/s]
Epoch [4/30], Train Loss: 1.5855, Val Loss: 1.8044, Val Acc: 0.5256, Top-3 Val Acc: 0.7394
100%|██████████| 167/167 [00:48<00:00, 3.48it/s]
Epoch [5/30], Train Loss: 1.4072, Val Loss: 1.8462, Val Acc: 0.5434, Top-3 Val Acc: 0.7576
100%|██████████| 167/167 [00:48<00:00, 3.42it/s]
Epoch [6/30], Train Loss: 1.2825, Val Loss: 1.5341, Val Acc: 0.5916, Top-3 Val Acc: 0.7795
100%|██████████| 167/167 [00:48<00:00, 3.48it/s]
Epoch [7/30], Train Loss: 1.1289, Val Loss: 1.4863, Val Acc: 0.6175, Top-3 Val Acc: 0.8215
100%|██████████| 167/167 [00:48<00:00, 3.42it/s]
Epoch [8/30], Train Loss: 0.9975, Val Loss: 1.4637, Val Acc: 0.6574, Top-3 Val Acc: 0.8465
100%|██████████| 167/167 [00:48<00:00, 3.42it/s]

```


Epoch [9/30], Train Loss: 0.9590, Val Loss: 1.2842, Val Acc: 0.7506, Top-3 Val Acc: 0.8715  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [10/30], Train Loss: 0.8947, Val Loss: 1.2675, Val Acc: 0.8055, Top-3 Val Acc: 0.8900  
 100%|██████████| 167/167 [00:48<00:00, 3.43it/s]  
 Epoch [11/30], Train Loss: 0.8137, Val Loss: 1.1348, Val Acc: 0.8452, Top-3 Val Acc: 0.8950  
 100%|██████████| 167/167 [00:48<00:00, 3.41it/s]  
 Epoch [12/30], Train Loss: 0.7621, Val Loss: 1.0010, Val Acc: 0.8699, Top-3 Val Acc: 0.9100  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [13/30], Train Loss: 0.6718, Val Loss: 0.9541, Val Acc: 0.8833, Top-3 Val Acc: 0.9214  
 100%|██████████| 167/167 [00:48<00:00, 3.49it/s]  
 Epoch [14/30], Train Loss: 0.6462, Val Loss: 0.8542, Val Acc: 0.8743, Top-3 Val Acc: 0.9320  
 100%|██████████| 167/167 [00:48<00:00, 3.49it/s]  
 Epoch [15/30], Train Loss: 0.6736, Val Loss: 0.8645, Val Acc: 0.8791, Top-3 Val Acc: 0.9390  
 100%|██████████| 167/167 [00:48<00:00, 3.48it/s]  
 Epoch [16/30], Train Loss: 0.6190, Val Loss: 0.8422, Val Acc: 0.8858, Top-3 Val Acc: 0.9416  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [17/30], Train Loss: 0.5844, Val Loss: 0.8021, Val Acc: 0.8938, Top-3 Val Acc: 0.9478  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [18/30], Train Loss: 0.3535, Val Loss: 0.7315, Val Acc: 0.9100, Top-3 Val Acc: 0.9500  
 100%|██████████| 167/167 [00:48<00:00, 3.39it/s]  
 Epoch [19/30], Train Loss: 0.2598, Val Loss: 0.7376, Val Acc: 0.9128, Top-3 Val Acc: 0.9517  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [20/30], Train Loss: 0.2242, Val Loss: 0.6816, Val Acc: 0.9171, Top-3 Val Acc: 0.9540  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [21/30], Train Loss: 0.2096, Val Loss: 0.6991, Val Acc: 0.9190, Top-3 Val Acc: 0.9555  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [22/30], Train Loss: 0.1983, Val Loss: 0.6416, Val Acc: 0.9210, Top-3 Val Acc: 0.9574  
 100%|██████████| 167/167 [00:48<00:00, 3.41it/s]  
 Epoch [23/30], Train Loss: 0.1980, Val Loss: 0.6611, Val Acc: 0.9227, Top-3 Val Acc: 0.9587  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [24/30], Train Loss: 0.2033, Val Loss: 0.6404, Val Acc: 0.9240, Top-3 Val Acc: 0.9600  
 100%|██████████| 167/167 [00:48<00:00, 3.43it/s]  
 Epoch [25/30], Train Loss: 0.1838, Val Loss: 0.6463, Val Acc: 0.9250, Top-3 Val Acc: 0.9610  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [26/30], Train Loss: 0.1757, Val Loss: 0.6486, Val Acc: 0.9266, Top-3 Val Acc: 0.9625  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [27/30], Train Loss: 0.1470, Val Loss: 0.6131, Val Acc: 0.9286, Top-3 Val Acc: 0.9632  
 100%|██████████| 167/167 [00:48<00:00, 3.49it/s]  
 Epoch [28/30], Train Loss: 0.1258, Val Loss: 0.6276, Val Acc: 0.9300, Top-3 Val Acc: 0.9645  
 100%|██████████| 167/167 [00:48<00:00, 3.42it/s]  
 Epoch [29/30], Train Loss: 0.1170, Val Loss: 0.5932, Val Acc: 0.9316, Top-3 Val Acc: 0.9655  
 100%|██████████| 167/167 [00:48<00:00, 3.47it/s]  
 Epoch [30/30], Train Loss: 0.1121, Val Loss: 0.5690, Val Acc: 0.9340, Top-3 Val Acc: 0.9667








Top-3 Accuracy vs. No. of epochs



Loss vs. No. of epochs



Adjusted Test set results: {'test\_loss': 0.4567, 'test\_acc': 0.9154, 'top3\_test\_acc': 0.9571}

+ Number of FLOPs: 0.60G